

Same AI. Different Results.

The title of this talk is Same AI, Different Results. I like it because it sounds almost unfair. We all have access to strong models now. We all have chat windows. We all have tools that can write, summarize, code, and plan. So why do the outcomes differ so much? Why does one team get something useful while another gets polished nonsense? My argument is simple. The difference is not mainly the model. The difference is whether the AI can work with the memory of the organization. In the age of AI, intelligence is becoming common. Context is not. Judgment is not. Shared memory is not. That is where the real difference starts.

NEXT

The Familiar Scene

To make that difference concrete, let me start with a scene most of us know. Two people ask almost the same thing. Maybe it is a teacher asking for a lesson plan. Maybe it is a manager asking for a brief. Maybe it is a developer asking for an implementation plan. Both answers sound good. Both are fluent. But one fits the local reality and the other does not. One uses the right language. The other ignores a rule, misses an exception, or suggests something the team already rejected. The prompt was similar. The model was similar. The hidden context was not. That is the basic puzzle behind this talk.

Why Better Prompting Is Not Enough

That familiar scene often makes us blame the prompt. We say we should have asked better. Sometimes that is true. But after a while that becomes a strange way to run a company. It means every person has to carry the organization in their own head and paste it into the chat again and again. Please remember our standards. Please remember our customer promises. Please remember why we built it this way. Please remember what legal said last month. That is not a system. That is humans acting as memory sticks for machines. The more agents we use, the more expensive that becomes. The memory needs a better home.

NEXT

AI Is A Shock. Coherence Is The Question.

The reason that memory needs a better home is that AI is not a small feature. It is a general-purpose technology, and I borrow this framing from Hans Kári Vang's talk at the AI strategy launch at Smæran: roughly 25 percent of work is taken over, 25 percent is amplified, and 50 percent is affected indirectly because expectations, handoffs, and coordination change. The exact numbers are less important than the shape: almost all work feels the shock, but in different ways. That is why the jobs debate can be too narrow. The better question for this room is this: when the tools multiply, what keeps the work coherent? What keeps a company, a public institution, or a professional practice from becoming disconnected outputs? That is where knowledge management becomes strategic again.

What Gets Lost

To see why coherence is hard, look at what knowledge work actually contains. On this slide I call the visible side the state clock: the document, the policy, the slide deck, the code, the answer to the customer. But underneath it sits the reasoning clock. Why did we choose this path? What risk were we trying to avoid? Why do we use this wording and not that one? Which exception matters in practice? An AI connected only to files can often see the state of the work. It cannot automatically see the reasoning that produced it. That is why generic AI can sound smart and still be locally wrong.

Two Kinds Of Context

That two-clock picture gives us a useful distinction. First, there is state context: what is true now. The current numbers. The current file. The current ticket. The current state of the world. That matters; agents need it to be accurate. But there is also reasoning context: why it became true. Why the team chose this path. What failed before. What constraint mattered. What exception keeps coming back. That is what most systems lose. I like the shorthand: data context makes agents accurate. Decision context makes agents institutional. Usable is strongest in the second half of that sentence.

NEXT

What Usable Is

The product answer to that second kind of context is Usable. So what is Usable really? Not mainly a knowledge base. Not a folder of PDFs. Not a wiki nobody updates. Not a frozen copy of every database. The best short description I found in our own material is this: Usable is organizational memory for AI agents. More precisely, governed organizational memory. A place where important decisions, standards, precedents, patterns, and lessons become available to people and agents at the moment work is happening. The point is not to make storage look clever. The point is to make real work less forgetful.

What Usable Is Not

The boundary is just as important as the memory itself. We should not pretend that all truth should live inside one memory system. That would create stale confidence very quickly. Live facts should often stay in the live systems where they belong. The current customer status. The current inventory. The current operational state. Usable is strongest around durable learning and decision context. Store the rules, the reasons, the standards, the lessons, the map of where truth lives, and the permissions around it. Then let the agent fetch volatile facts live. That is a much cleaner product story.

NEXT

The Loop That Compounds

When you combine governed memory with live systems, the work loop changes. Without a memory layer, the loop is simple and wasteful: ask, get answer, forget. The next person starts again. The next agent starts again. With a memory layer, the loop becomes read before acting, act, capture what mattered, govern it, and reuse it next time. That sounds small, but it changes everything. It moves AI from disposable help to compounding capability. It also reflects the product pattern we keep seeing: the agent is grounded in a workspace, it can reach the right local tools, and when the work is done it leaves something useful behind.

NEXT

Proof 1: TAKS

A public example makes that loop less abstract: TAKS. The assistant is called Eyðei, which sounds like AI in Faroese, and TAKS even gave the wider cast names like Tollakur and Skatta Ró. It is friendly on the surface, but the real point is deeper. Tax is high-trust work. People ask about deadlines, customs, deductions, and consequences. A generic bot is not enough. The assistant has to be grounded in official material. It has to be careful with sources. It has to improve from what happens in the real workflow. That is why this story matters. The value is not the chat face. The value is the governed memory behind it.

NEXT

Proof 2: Julius

At Thursday's Usable Fragments event at Reinsariid, Julius gave the software-development version of the same problem. His point was not just that AI makes developers faster. His point was that speed creates a new process problem. Three good developers can ask the same AI for the same feature and get three reasonable but different implementations. None are obviously crazy. They just do not fit together. Then the team pays the alignment tax at review time. Julius showed a better loop: start from a PRD, split into aligned tasks, enhance the plan with standards and prior learning before code is written, then write the lesson back after the commit. That moves alignment from review time to plan time.

NEXT

The Pattern Repeats

The software story was not an isolated case. Across the Usable Fragments talks, the same missing layer showed up in very different domains. In Jaspur's strategy work, the project memory became part of the deliverable. The value was not only the final document, but the reasoning trail behind it. In Jóhann's game-development work, the breakthrough came when the project stopped living only in the latest prompt. The agent got better when it could return to a workspace with history. Different domain. Same pattern. The tool gets more useful when it can stand on accumulated local context instead of improvising from scratch every time.

More Than One Interface

Brian and Ajit took the same pattern one step further. Brian's operators showed that the important thing is not the surface. Slack, IDE, browser, terminal, phone: those are interfaces. The durable part is the memory layer underneath, plus the harness around it. Ajit then gave the conceptual language for the frontier. If we want agents to improve, we need more than retrieval. We need traces of how decisions were made and what happened next. That is when memory starts becoming precedent. And that is when knowledge management begins to feel modern again instead of old-fashioned.

NEXT

What A Decision Trace Is

Let me slow down on the phrase decision trace, because it is easier to explain than it sounds. It says: what did we observe, what constraint mattered, what tradeoff did we make, what action did we take, and what happened afterwards? That is different from a log. A log says what happened. A decision trace tries to preserve why it happened. This is one of the deepest points in the current Usable story. Organizations already store lots of output. Much fewer store chains of judgment that an agent can actually reuse. Once you preserve that chain, memory becomes more valuable than search alone.

Precedent And Boundaries

Once you see decisions as reusable traces, Ajit's context-graph framing becomes useful. The goal is not just to find similar text. The goal is to understand relationships between decisions, exceptions, actions, and outcomes. But there is another deep point here. One big company brain is not the answer. The hard part is governing the overlap. The sales assistant should not see board material. The engineering agent should not get every private HR note. Good organizational memory is not only rich. It is well-bounded. That is why governed memory matters so much.

NEXT

Every Profession Has This

The reason I keep saying profession, not software team, is that every knowledge profession has local judgment that generic AI cannot guess. A teacher has examples that worked with a certain group. A lawyer has clause history and negotiation rationale. An engineer has incidents, standards, and scars. A manager has tradeoffs, promises, and stakeholder history. A researcher has sources, gaps, and caveats. A founder has bets, signals, and lessons learned the hard way. The forms change, but the pattern does not. Good work depends on more than facts. It depends on remembered judgment in a specific context.

Personal Memory, Team Memory

That professional memory exists at two scales: personal memory and team memory. Your own best examples, checklists, and explanations matter. So do the standards that live only in the heads of experienced teammates. The opportunity is to make that knowledge reusable without pretending everyone should see everything. This is why I keep coming back to the word governed. The future is not one giant pile of memory. It is shared memory where it helps, local memory where it belongs, and clear boundaries in between. That is a more realistic vision of AI at work.

NEXT

Monday Morning, Part One

If the idea is clear, the practical question is: what should someone in this room do on Monday morning? First, pick one repeated workflow. Not the whole organization. One workflow. Second, capture the decisions and standards that shape good work in that workflow. What do experienced people know that newcomers or agents do not? Third, connect the agent to that memory before it acts. Do not wait until the review stage to add context. Bring the context earlier, into the plan and the first draft. That alone changes the quality of what comes back.

Monday Morning, Part Two

Once that first workflow has memory before the work starts, add two practical habits. Let the agent reach the right local tools instead of guessing around them. And make the end of the task include a memory update. What changed? What was learned? What should the next person or next agent know? That is enough to start. You do not need a giant transformation program. You need one loop that stops throwing away learning. If you do that well, the benefit is not only speed. It is calm. Fewer repeated mistakes. Less re-explaining. Better first drafts. Better second runs.

NEXT

Close

That brings me back to the title: Same AI. Different Results. In the age of AI, many people can buy access to intelligence. That means advantage moves somewhere else. It moves to owned memory. It moves to judgment made durable. It moves to whether your agents can work with the reality of your world instead of only the patterns of the internet. That is the story I think Usable is really telling now. Not just a knowledge base. Not just search. Governed organizational memory for AI agents. Build that layer well and the same model starts giving you very different outcomes. Thank you.